

Table 1: Post-hoc validation-only IPID-space sensitivity. Values are equally weighted macro means over eight primary cells (six matched traces per cell). The difference is 65,536 minus 2,048; intervals use 5,000 paired-trace bootstrap resamples stratified by cell. B5 uses the unchanged locked thresholds. Rates and J are on the 0–1 scale; these are descriptive intervals, not a new held-out evaluation.

Metric	2,048	65,536	Difference	95% CI
Entropy AP	0.5780	0.5041	-0.0739	[-0.0786, -0.0696]
Unique-ratio AP	0.7512	0.5196	-0.2315	[-0.2405, -0.2224]
B5 attack alert	0.5911	0.5942	+0.0031	[+0.0015, +0.0047]
B5 benign trigger	0.5814	0.5936	+0.0122	[+0.0103, +0.0144]
B5 FNR	0.4089	0.4058	-0.0031	[-0.0047, -0.0015]
B5 J	0.0097	0.0006	-0.0092	[-0.0121, -0.0067]

Table 2: Offline orphan-fragment baselines on the no-enforcement runtime traces. Load denotes background non-initial fragments/s. Both criteria also require $n \geq 8$ in a two-second window. Alert rates count any activation from query start to completion; each traffic-rate cell has 20 recreated-stack runs and 1,000 trials. The orphan ratio O uses causal first-fragment matching by source, destination, IP protocol and IPID. No orphan policy was enforced, so these are detection rates, not poisoning-prevention rates.

Load (/s)	Criterion	Benign alert (%)	Attack alert (%)	J
12	$O > 0$	0.00	0.00	0.0000
12	$O \geq 0.5$	0.00	0.00	0.0000
200	$O > 0$	0.00	0.00	0.0000
200	$O \geq 0.5$	0.00	0.00	0.0000