

Table 1: Runtime outcomes in the original campaign and the separately analyzed 2×2 follow-up. Each row contains 20 recreated-stack runs with 50 trials per run. Malicious-answer and no-answer rates are trial percentages. Affected runs contain at least one malicious answer; intervals are exact 95% Clopper–Pearson intervals for the affected-run percentage. Fixed/Flood denote the original matched fixed-IPID attack/sweep flood. In the follow-up, B/A denote benign/attack and 12/200 denote background non-initial fragments/s under the same IPID sweep.

Policy	Workload	Malicious (%)	No answer (%)	Affected runs	95% CI (%)
<i>Original campaign</i>					
B0	Fixed	97.5	0.9	20/20	83.2–100.0
B0	Flood	99.5	0.5	20/20	83.2–100.0
B1+TC	Fixed	0.0	0.2	0/20	0.0–16.8
B1+TC	Flood	0.0	0.8	0/20	0.0–16.8
Locked B5+TC	Fixed	98.0	0.4	20/20	83.2–100.0
Locked B5+TC	Flood	0.0	0.6	0/20	0.0–16.8
Native drop	Fixed	0.0	100.0	0/20	0.0–16.8
Native drop	Flood	0.0	100.0	0/20	0.0–16.8
<i>Post-hoc-designed 2×2 follow-up</i>					
B0	B12	0.0	0.2	0/20	0.0–16.8
B0	B200	0.0	0.1	0/20	0.0–16.8
B0	A12	100.0	0.0	20/20	83.2–100.0
B0	A200	99.6	0.2	20/20	83.2–100.0
B1+TC	B12	0.0	0.4	0/20	0.0–16.8
B1+TC	B200	0.0	0.5	0/20	0.0–16.8
B1+TC	A12	0.0	0.2	0/20	0.0–16.8
B1+TC	A200	0.0	0.5	0/20	0.0–16.8
Locked B5+TC	B12	0.0	0.3	0/20	0.0–16.8
Locked B5+TC	B200	0.0	0.7	0/20	0.0–16.8
Locked B5+TC	A12	99.8	0.2	20/20	83.2–100.0
Locked B5+TC	A200	0.0	0.2	0/20	0.0–16.8
Native drop	B12	0.0	100.0	0/20	0.0–16.8
Native drop	B200	0.0	100.0	0/20	0.0–16.8
Native drop	A12	0.0	100.0	0/20	0.0–16.8
Native drop	A200	0.0	100.0	0/20	0.0–16.8