

Table 1: Benign runtime costs in the original runtime campaign. Rates are percentages; latency is median and 95th percentile (ms) of 1,000 query attempts per cell, including timeouts and excluding cache probes. These quantiles are descriptive; 20 recreated-stack runs are the independent units. Native dropping uses concurrent queries; the other policies use sequential queries.

Policy	Load	TC %	TCP %	Med. ms	p95 ms	None %
B0	2.5/s	0.0	0.0	334.3	361.2	0.8
B0	12/s	0.0	0.0	336.9	363.2	0.4
B1+TC	2.5/s	100.0	100.0	376.0	409.7	0.7
B1+TC	12/s	100.0	100.0	380.5	420.2	1.3
B5+TC	2.5/s	0.0	0.0	334.2	363.4	0.7
B5+TC	12/s	0.0	0.0	337.0	365.0	0.3
Native drop	2.5/s	0.0	0.0	2010.5	2022.1	100.0
Native drop	12/s	0.0	0.0	2010.7	2021.7	100.0

Table 2: Benign runtime costs in the diverse-IPID factorial follow-up. Rates are percentages; latency is median and 95th percentile (ms) of 1,000 query attempts per cell, including timeouts and excluding cache probes. These quantiles are descriptive; 20 recreated-stack runs are the independent units. Native dropping uses concurrent queries; the other policies use sequential queries.

Policy	Load	TC %	TCP %	Med. ms	p95 ms	None %
B0	12/s	0.0	0.0	337.8	363.2	0.2
B0	200/s	0.0	0.0	340.8	367.3	0.1
B1+TC	12/s	100.0	100.0	378.6	412.0	0.4
B1+TC	200/s	100.0	100.0	399.0	428.1	0.5
B5+TC	12/s	0.0	0.0	338.1	364.1	0.3
B5+TC	200/s	100.0	100.0	398.6	434.2	0.7
Native drop	12/s	0.0	0.0	2010.2	2021.3	100.0
Native drop	200/s	0.0	0.0	2011.1	2021.6	100.0

Table 3: Original runtime outcomes, replacing the former runtime figure. Malicious-answer and no-answer rates are trial percentages (1,000 trials per cell); affected runs contain at least one malicious answer. Intervals are exact 95% Clopper–Pearson intervals for the affected-run percentage; 20 recreated-stack runs per cell. Cache insertion was not independently verified in this campaign.

Policy	Attack	Mal. %	None %	Runs	95% CI
B0	Fixed	97.5	0.9	20/20	83.2–100.0
B0	Flood	99.5	0.5	20/20	83.2–100.0
B1+TC	Fixed	0.0	0.2	0/20	0.0–16.8
B1+TC	Flood	0.0	0.8	0/20	0.0–16.8
B5+TC	Fixed	98.0	0.4	20/20	83.2–100.0
B5+TC	Flood	0.0	0.6	0/20	0.0–16.8
Native drop	Fixed	0.0	100.0	0/20	0.0–16.8
Native drop	Flood	0.0	100.0	0/20	0.0–16.8

Table 4: Locked-B5 mechanism rates in the original campaign, in percent. Each cell contains 20 recreated-stack runs and 1,000 trials. Each of the four run-level rates is constant within a workload, so its percentile bootstrap 95% interval equals its estimate.

Attack	Trigger %	Drop %	TC %	TCP %
Fixed	0.0	0.0	0.0	0.0
Flood	100.0	100.0	100.0	100.0

Table 5: Diverse-IPID factorial follow-up: forged-tail injection on/off crossed with background load (12 or 200 non-initial fragments/s). B5 state is the reconstructed locked-B5 state at query time, including passive state under other policies. All outcome columns are trial percentages; 20 runs and 1,000 trials per cell. Cache denotes a malicious A record present at the immediate post-query probe, not long-term persistence. This campaign is analyzed separately from the original runtime experiment.

Policy	Traffic /s	B5 state %	Mal. %	Cache %	None %
B0	Benign 12	0.0	0.0	0.0	0.2
B0	Benign 200	100.0	0.0	0.0	0.1
B0	Attack 12	0.0	100.0	100.0	0.0
B0	Attack 200	100.0	99.6	99.8	0.2
B1+TC	Benign 12	0.0	0.0	0.0	0.4
B1+TC	Benign 200	100.0	0.0	0.0	0.5
B1+TC	Attack 12	0.0	0.0	0.0	0.2
B1+TC	Attack 200	100.0	0.0	0.0	0.5
B5+TC	Benign 12	0.0	0.0	0.0	0.3
B5+TC	Benign 200	100.0	0.0	0.0	0.7
B5+TC	Attack 12	0.0	99.8	100.0	0.2
B5+TC	Attack 200	100.0	0.0	0.0	0.2
Native drop	Benign 12	0.0	0.0	0.0	100.0
Native drop	Benign 200	100.0	0.0	0.0	100.0
Native drop	Attack 12	0.0	0.0	0.0	100.0
Native drop	Attack 200	100.0	0.0	0.0	100.0